

Project 1

Yash Prashant Gokhale

College of Professional Studies, Northeastern University

ALY6000: Introduction to Analytics

Prof. Richard He

September 25, 2023

Introduction

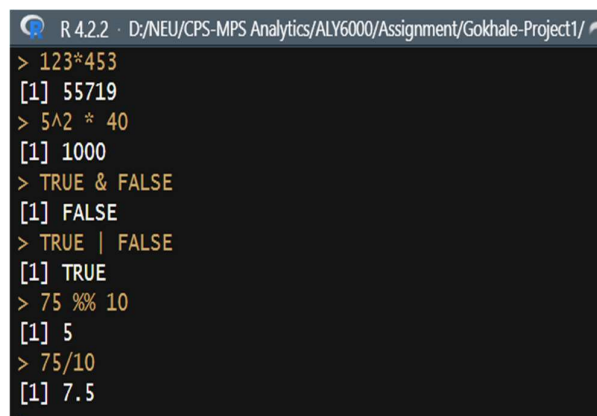
This project acts as an introductory session to the R language and at the same time gives an idea as to how statistics can be seamlessly performed through the language. Vectors and Data Frames are the data types that have been stressed throughout the project.

Various handy, in-built functions such as “c”, “seq”, “rep”, statistical functions like sum, mean, median, max, min are used. Accessing data from a csv file using “data_csv()” function and then using functions like head(), name(), select(), arrange(), filter(), slice(), mutate() are implemented so as to have a clearer understanding of what the data frame is, which is almost always the very first step of EDA.

Function like the ggplot() is used to visualize a part of data so as to convey the story behind the data graphically.

Key Findings

- (1) Along with the use of sophisticated statistical functions, R can also be used to compute simple calculations as shown below.



```
R 4.2.2 · D:/NEU/CPS-MPS Analytics/ALY6000/Assignment/Gokhale-Project1/
> 123*453
[1] 55719
> 5^2 * 40
[1] 1000
> TRUE & FALSE
[1] FALSE
> TRUE | FALSE
[1] TRUE
> 75 %% 10
[1] 5
> 75/10
[1] 7.5
```

- (2) The code snippet below shows as to how the Broadcasting works in R. The first line adds 20 to all the elements of the vector ‘second_vector’ while the next line multiplies all the elements of the ‘second_vector’ by 20. The next two lines compare all the elements of the vector ‘second_vector’ with 20 and accordingly return a Boolean vector with TRUE corresponding to 1 and FALSE corresponding to 0.

These lines in no way mutate or change the original vector ‘second_vector’.

```
Console Terminal Background Jobs
R 4.2.2 D:/NEU/CPS-MPS Analytics/ALY6000/Assignment/Gokhale-Project1/
> # 20 has been added to all the elements of the vector
> second_vector+20
[1] 30 32 34 36 38 40 42 44 46 48 50
>
> # All the elements in the vector get multiplied by 20
> second_vector*20
[1] 200 240 280 320 360 400 440 480 520 560 600
>
> # Returns the elements in the given vector that are greater than 20 as TRUE and remaining once as FALSE.
> second_vector >= 20
[1] FALSE FALSE FALSE FALSE FALSE TRUE TRUE TRUE TRUE TRUE TRUE
>
> # Returns the elements in the given vector that are not equal to 20 as TRUE and the remaining as FALSE.
> second_vector != 20
[1] TRUE TRUE TRUE TRUE TRUE FALSE TRUE TRUE TRUE TRUE TRUE
```

- (3) seq() function is used to generate a sequence of elements for a vector. The parameters like 'from' and 'to' act as starting and ending points for the sequence generation while 'by' can be used to provide an equal increment/decrement.

```
> reverse_numbers <- seq(from = 100, to = -100, by = -3)
> reverse_numbers
[1] 100 97 94 91 88 85 82 79 76 73 70 67 64 61 58 55 52 49 46 43
[21] 40 37 34 31 28 25 22 19 16 13 10 7 4 1 -2 -5 -8 -11 -14 -17
[41] -20 -23 -26 -29 -32 -35 -38 -41 -44 -47 -50 -53 -56 -59 -62 -65 -68 -71 -74 -77
[61] -80 -83 -86 -89 -92 -95 -98
```

- (4) A Boolean vector is created using the combine function and 'first_vector' as its parameter. The corresponding 'TRUE' values are returned.

```
> # The second and the Third element from the first_vector were returned.
> vector_from_boolean_brackets <- first_vector[c(FALSE,TRUE,FALSE, TRUE)]
> vector_from_boolean_brackets
[1] 12 5
```

- (5) "second_vector >= 20" returns a Boolean vector which states 'TRUE' for the values that are above 20 and 'FALSE' for the values less than 20.

```
> # Returns the elements from the second_vector which are above 20 as TRUE while others as FALSE.
> second_vector >= 20
[1] FALSE FALSE FALSE FALSE FALSE TRUE TRUE TRUE TRUE TRUE TRUE
```

- (6) 'ages_vector' returns the values that are above 20 using Booleans as a mode of filtrations.

```
> # Returns the values which are greater then 20 from the 'ages_vector'
> ages_vector [ages_vector>=20]
[1] 20 22 24 26 28 30
```

- (7) By passing negative indexes to the combine function, we can drop the corresponding elements from the vector as displayed below.

```
R 4.2.2 · D:/NEU/CPS-MPS Analytics/ALY6000/Assignment/Gokhale-Project1/ ↗
> middle_grades_removed <- grades[c(-3,-4)]
> middle_grades_removed
[1] 96 100 81 72
```


- (8) `set.seed()` is used to create the exact same random variables every time while the `runif()` is used to create a random variables with the arguments as minimum(starting point), maximum(end point) and the actual length of the vector.

```
> # set.seed() is used to create the exact same Random Variables every time.
> # runif() is used to create the vector of given length, which each value being random.
> set.seed(5)
> random_vector <- runif(n=10, min = 0, max = 1000)
> random_vector
[1] 200.2145 685.2186 916.8758 284.3995 104.6501 701.0575 527.9600 807.9352 956.5001
[10] 110.4530
```

- (9) The below code shows how to create random variables using normal distribution. Here a 1000 values with a mean value of 50 and the standard deviation of 15 are sampled out from a Gaussian Distribution using the “`rnorm`” function.

```
Console Terminal Background Jobs
R 4.2.2 · D:/NEU/CPS-MPS Analytics/ALY6000/Assignment/Gokhale-Project1/
> set.seed(5)
> random_vector <- rnorm(n=1000, mean = 50, sd = 15)
> random_vector
[1] 37.387168 70.765390 31.167622 51.052141 75.671613 40.956380 42.917504
[8] 40.469430 45.713395 52.071623 68.414455 37.973308 33.794111 47.636985
[15] 33.923599 47.915208 41.040304 17.240499 53.612259 46.109669 63.507679
[22] 64.128041 72.019429 60.601416 62.285134 45.597772 71.278836 72.481607
[29] 40.143769 37.208068 54.738726 66.645413 83.231909 68.256555 72.188327
[36] 64.273607 34.857010 19.992909 23.567212 47.860878 73.250906 37.963652
[43] 48.881316 78.435019 43.151466 58.433350 36.694872 43.096331 39.135073
[50] 48.961833 71.948728 52.815891 65.330343 41.122478 48.316990 36.125704
[57] 61.299572 48.310864 49.038636 53.499129 32.951258 62.822456 41.324444
[64] 57.445423 38.599131 44.879206 18.465063 45.474466 30.914248 45.805008
[71] 46.938540 46.615787 55.205427 50.485518 56.202969 47.669773 64.602281
[78] 51.816352 52.837605 41.556724 57.476242 23.865463 64.632936 49.638757
[85] 60.135267 39.345356 85.808490 42.898520 48.863412 42.172399 63.890707
[92] 34.063832 58.355508 63.510959 64.849185 55.754121 44.801243 41.897161
[99] 47.261666 49.110505 20.069195 67.029669 60.136918 53.127249 49.132315
[106] 63.407171 46.567019 20.515210 38.697343 69.202274 35.706426 74.335691
[113] 89.002130 52.094728 29.739205 61.983965 26.675062 56.955801 50.786443
[120] 46.969523 67.562846 63.272673 30.231671 25.351236 65.888756 54.351254
[127] 43.999498 68.646437 29.503842 28.378800 70.228236 20.322075 31.385741
[134] 48.439413 60.994594 56.835194 54.321193 33.894636 59.731138 54.487434
[141] 38.060075 49.559699 82.703536 64.361277 45.424270 43.723950 51.499311
[148] 46.552856 28.771777 44.111017 64.191328 61.276563 42.239347 62.125040
[155] 40.781972 68.573884 44.928573 67.945495 43.350224 52.791723 10.679828
[162] 83.693819 51.401475 74.409201 42.336237 40.109287 49.397148 48.219590
[169] 49.705147 42.714823 28.397787 52.156533 31.481200 23.712482 49.467556
[176] 54.980524 73.584324 33.957941 63.744298 41.075107 82.724700 39.743401
[183] 61.250888 64.615740 31.032898 45.838679 47.159020 44.239626 61.108820
[190] 32.474924 60.013080 55.493554 42.275855 56.758524 47.184194 70.086041
[197] 62.243288 51.233026 40.237059 60.896135 48.294827 45.573488 64.837527
[204] 38.373023 54.138474 56.161725 59.167747 64.048561 44.486874 61.105651
[211] 68.277996 59.437017 57.916195 42.916170 62.355727 43.583176 47.860341
[218] 71.281746 57.307009 59.051622 53.162493 49.500512 80.377955 44.438199
```

(10) R can be extensively used to create engaging visualizations, like histogram as shown below.
The hist function is passed with “random_vector” to plot the histogram with the 1000 samples with the mean of 50 and SD of 15.



(11) The following lines of R code prove as to how flexible and versatile the R language is when it comes to analyzing the data frames. The readily available functions like head, select, arrange, filter, mutate and slice are extremely useful to get an overview as to how the data is structured and do the data conversions to gain insights out of it.

```
first_dataframe <- read_csv("ds_salaries.csv")

head(first_dataframe)
head(first_dataframe, n=7)
names(first_dataframe)

smaller_dataframe <- select(first_dataframe, job_title, salary_in_usd)
smaller_dataframe

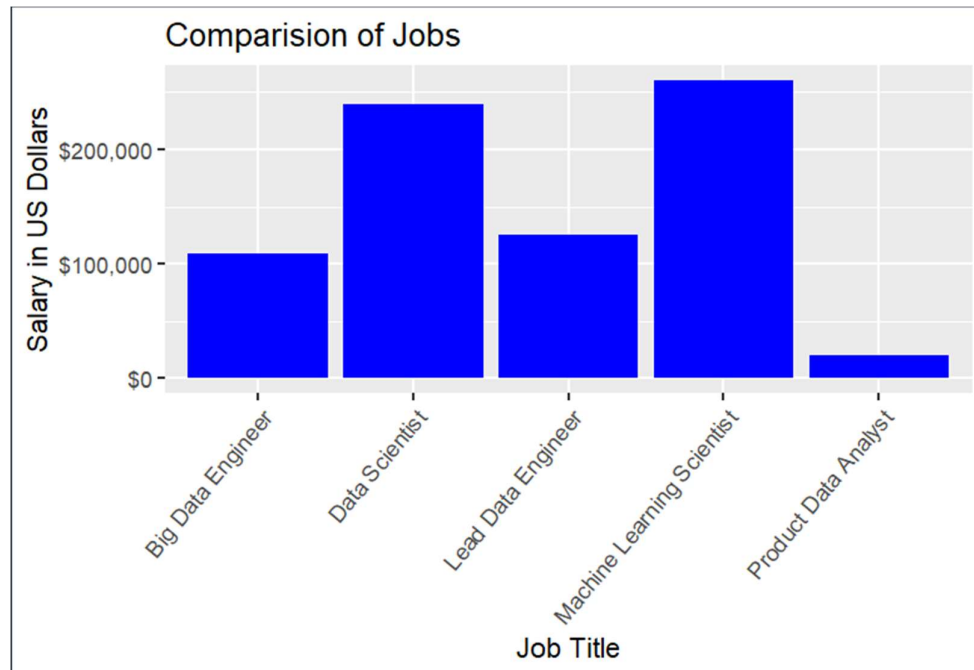
better_smaller_dataframe <- arrange(smaller_dataframe, desc(salary_in_usd))
better_smaller_dataframe
better_smaller_dataframe <- filter(smaller_dataframe, salary_in_usd > 80000)
better_smaller_dataframe

better_smaller_dataframe <- mutate(smaller_dataframe, salary_in_euros = salary_in_usd / 1.08)
better_smaller_dataframe

better_smaller_dataframe <- slice(smaller_dataframe, 1, 1, 2, 3, 4, 10, 1)
better_smaller_dataframe

ggplot(better_smaller_dataframe) +
  geom_col(mapping = aes(x = job_title, y = salary_in_usd), fill = "blue") +
  xlab("Job Title") +
  ylab("Salary in US Dollars") +
  labs(title = "Comparison of Jobs") +
  scale_y_continuous(labels = scales::dollar) +
  theme(axis.text.x = element_text(angle = 50, hjust = 1))
```

- (12) As mentioned previously R has extensive functions and ability to create visualization which truly convey insights in a compelling graphical manner. ggplot() from tidyverse package is extremely useful to create visualizations from the dataset or data frames. Various parameters like data-frame, labels, title, legend, color can be passed to create graphs. In the assignment a bar graph showing the “Comparison of the Job” is visualized with x-axis being “Job Title” while the y-axis being “Salary in USD”



Conclusion

The assignment acted as a good exercise to gently introduce the R scripting language. It was quite good to learn and understand functions of R which have tremendous potential to not just to EDA but also carry out statistical analysis. R language also has a plethora of functions to visualize the data in a compelling manner. Overall, the assignment successfully helped me understand the basic functionality of R.

References

<https://www.rdocumentation.org/packages/pacman/versions/0.5.1>

<https://www.rdocumentation.org/packages/ggplot2/versions/0.9.0/topics/ggplot>

<https://www.rdocumentation.org/packages/tidyverse/versions/2.0.0>